



University
of
St Andrews



Parallel Functional Programming

Kevin Hammond
University of St Andrews, Scotland

<http://www.paraphrase-ict.eu>

<http://www.project-advance.eu>

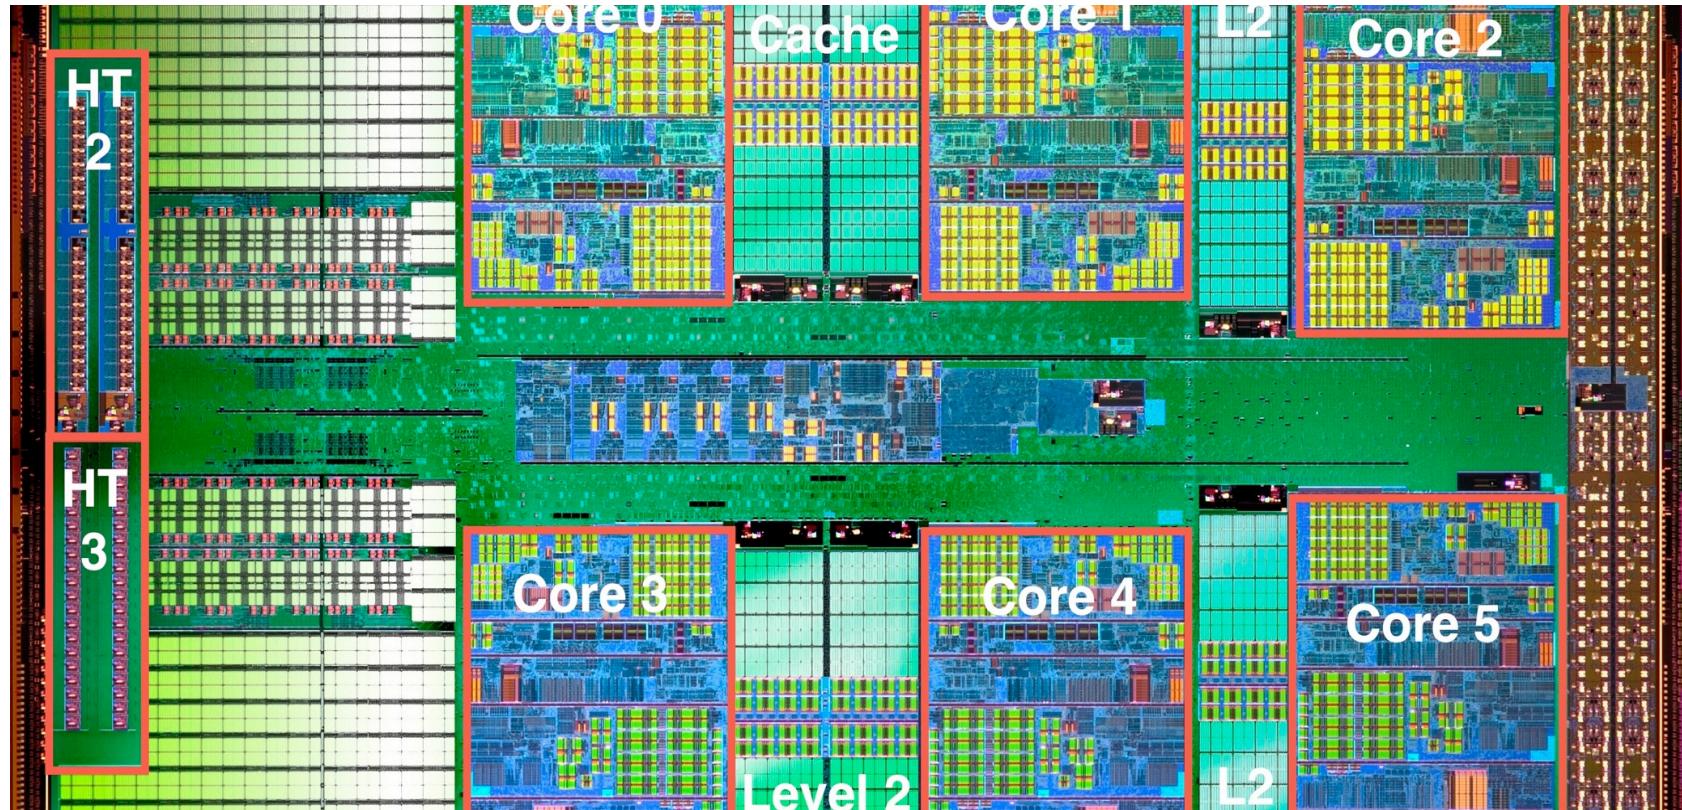
@paraphrase_fp7



PARAPHRASE



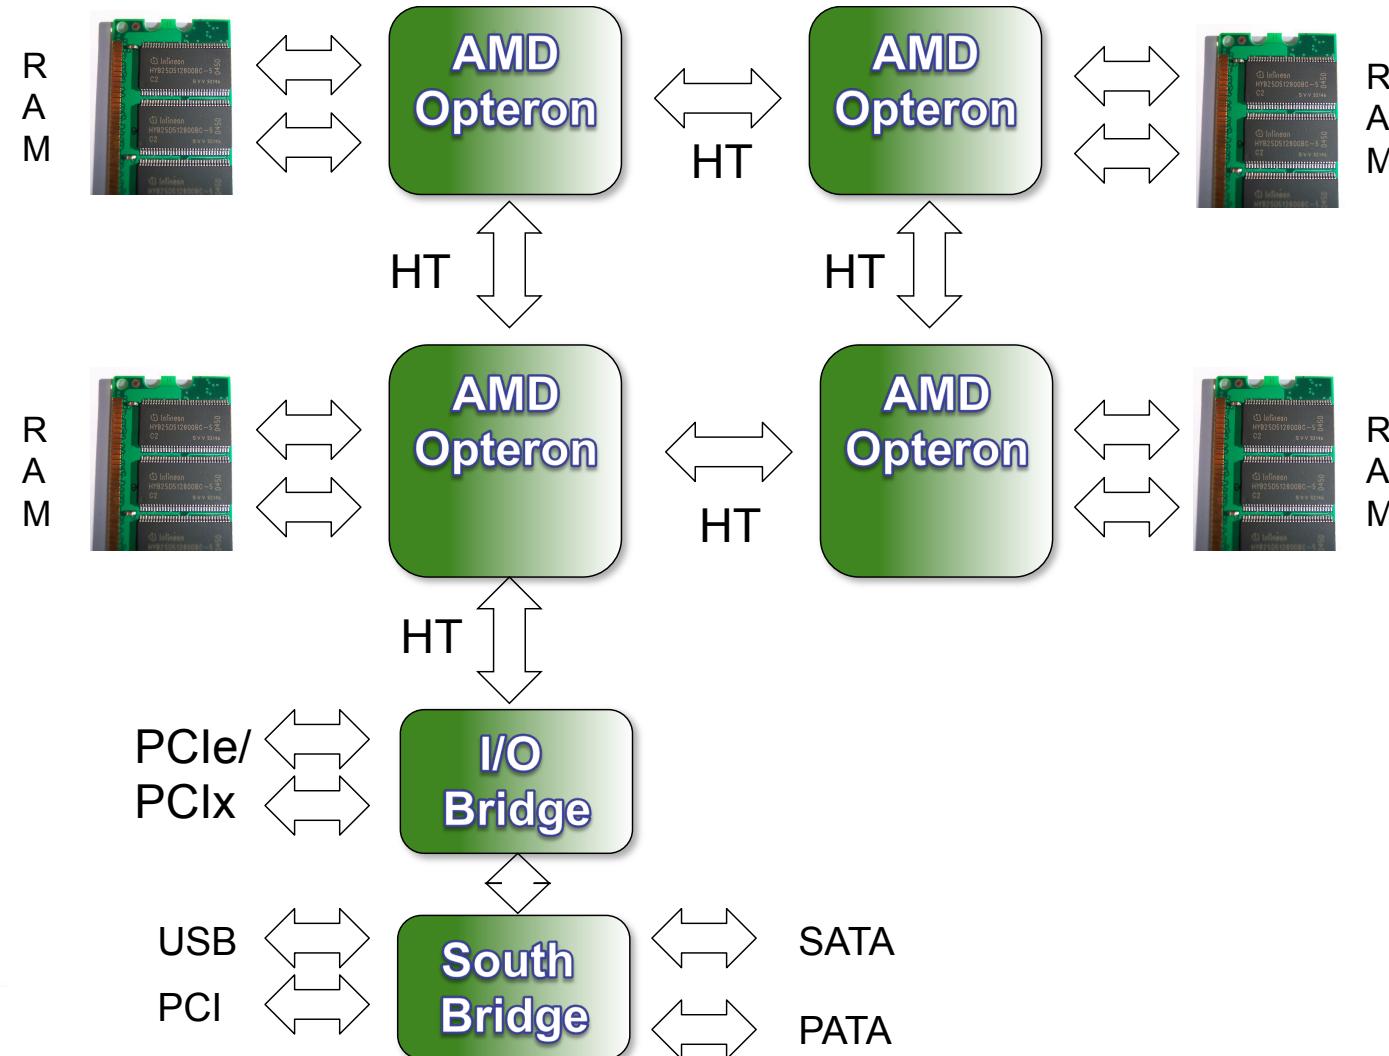
The Dawn of the Multicore Age



AMD Opteron Magny-Cours , 6-Core (source: wikipedia)

PARAPHRASE

The Near Future: Scaling toward Manycore



The Challenge

“Ultimately, developers should start thinking about *tens, hundreds, and thousands* of cores *now* in their algorithmic development and deployment pipeline.”

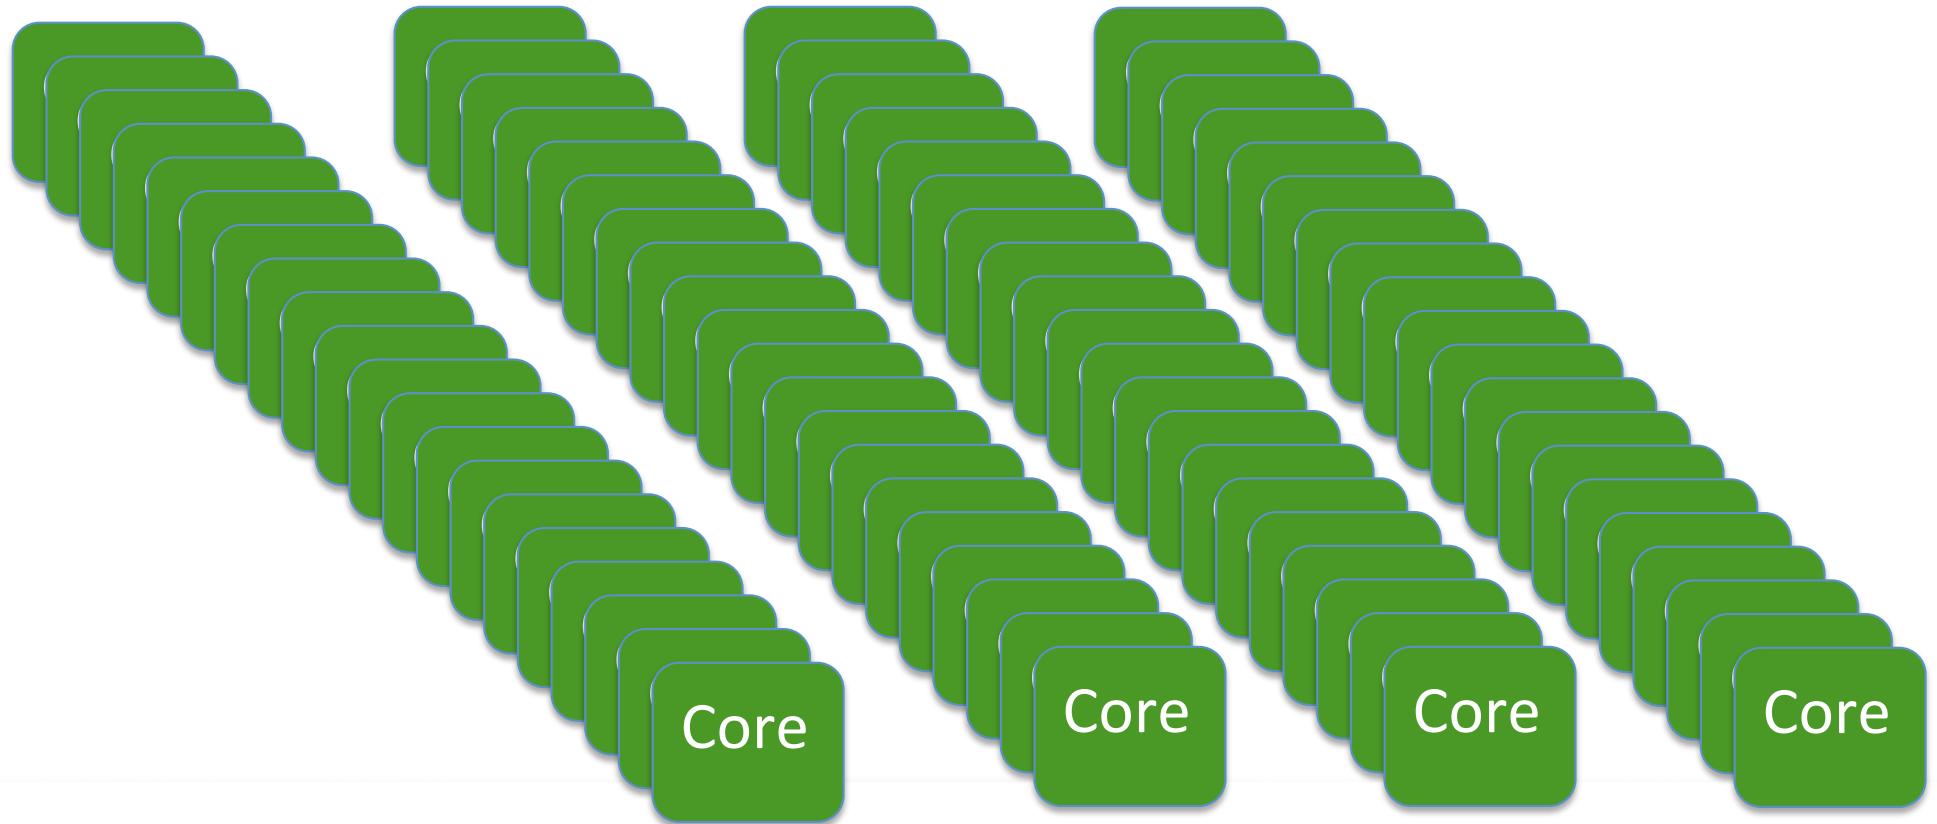
Anwar Ghuloum, Principal Engineer, Intel Microprocessor Technology Lab

“The dilemma is that a *large percentage* of mission-critical enterprise applications will not ``automagically'' run faster on multi-core servers. *In fact, many will actually run slower.* We must make it as easy as possible for applications programmers to exploit the latest developments in multi-core/many-core architectures, while still making it easy to target future (and perhaps unanticipated) hardware developments.”

**Patrick Leonard, Vice President for Product Development
Rogue Wave Software**

The Future: “megacore” computers?

- *Hundreds of thousands, or millions, of cores*



What will “megacore” computers look like?



- Probably *not* just scaled versions of today’s multicore
 - Perhaps hundreds of dedicated lightweight integer units
 - Hundreds of floating point units (enhanced GPU designs)
 - A *few* heavyweight general-purpose cores
 - Some specialised units for graphics, authentication, network etc
 - possibly *soft* cores (FPGAs etc)
 - *Highly heterogeneous*
- Probably *not* uniform shared memory
 - NUMA is likely, even hardware distributed shared memory
 - or even message-passing systems on a chip
 - *shared-memory will not be a good abstraction*



Laki (NEC Nehalem Cluster) and hermit (XE6)

Laki

- ▶ 700 dual socket Xeon 5560 2,8GHz (“Gainestown”)
- ▶ 12 GB DDR3 RAM / node
- ▶ Infiniband (QDR)
- ▶ 32 nodes with additional Nvidia Tesla S1070
- ▶ Scientific Linux 6.0

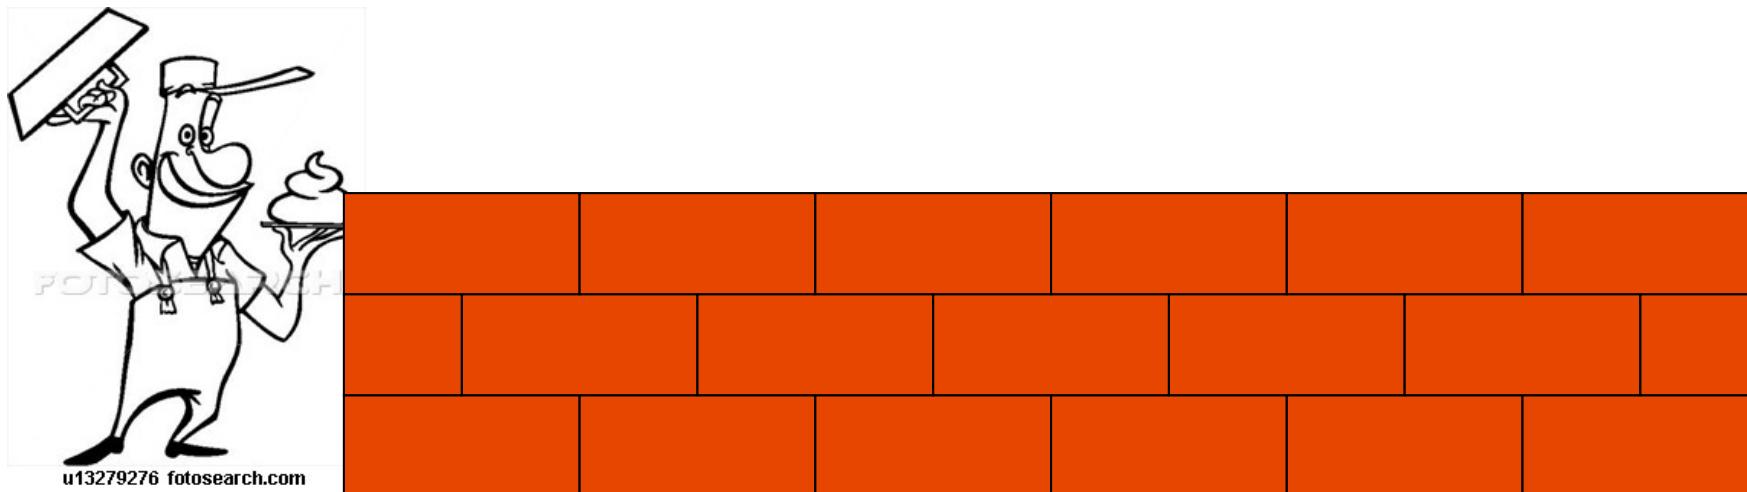
hermit (phase 1 step 1)

- ▶ 38 racks with 96 nodes each
- ▶ 96 service nodes and 3552 compute nodes
- ▶ Each compute node will have 2 sockets AMD Interlagos @ 2.3GHz 16 Cores each leading to 113.664 cores
- ▶ Nodes with 32GB and 64GB memory reflecting different user needs
- ▶ 2.7PB storage capacity @ 150GB/s IO bandwidth
- ▶ External Access Nodes, Pre- & Postprocessing Nodes, Remote Visualization Nodes

All future programming will be parallel

- No future system will be single-core
 - you must use parallel programming
- It's not just about performance
 - it's also about energy usage
- If we don't solve the multicore challenge, then all other CS advances won't matter!
 - user interfaces
 - cyber-physical systems
 - robotics
 - games
 - ...

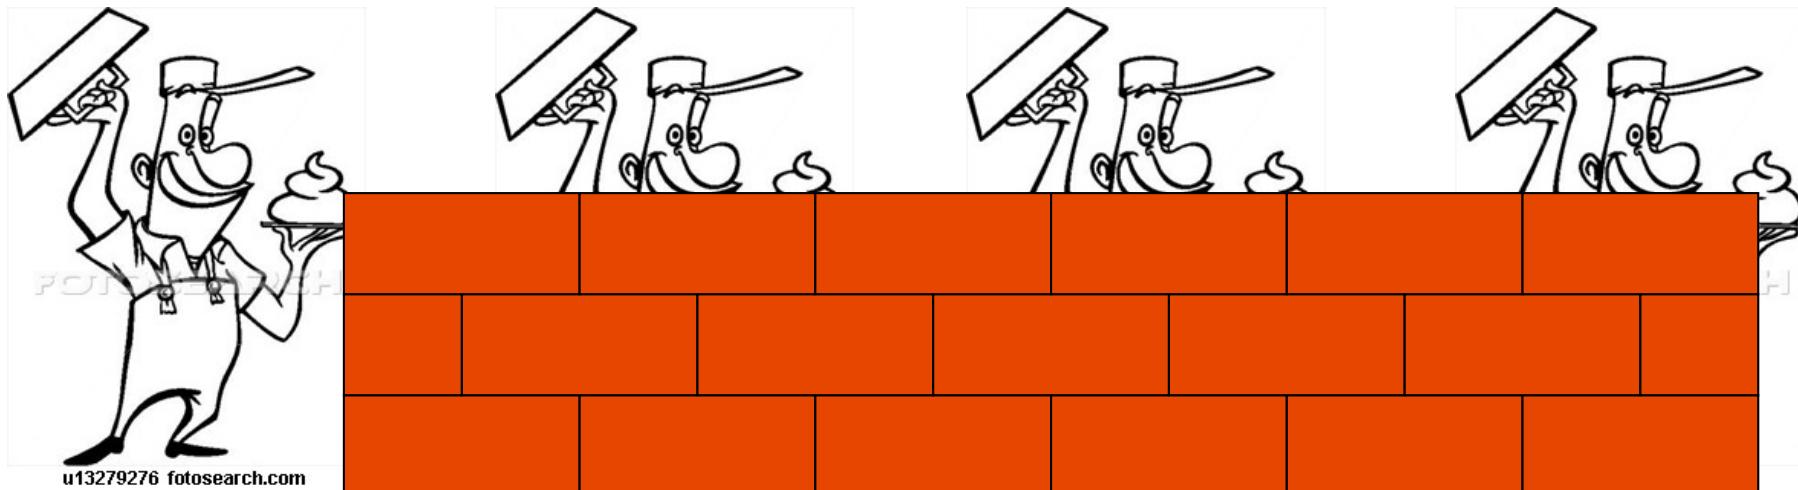
How to build a wall



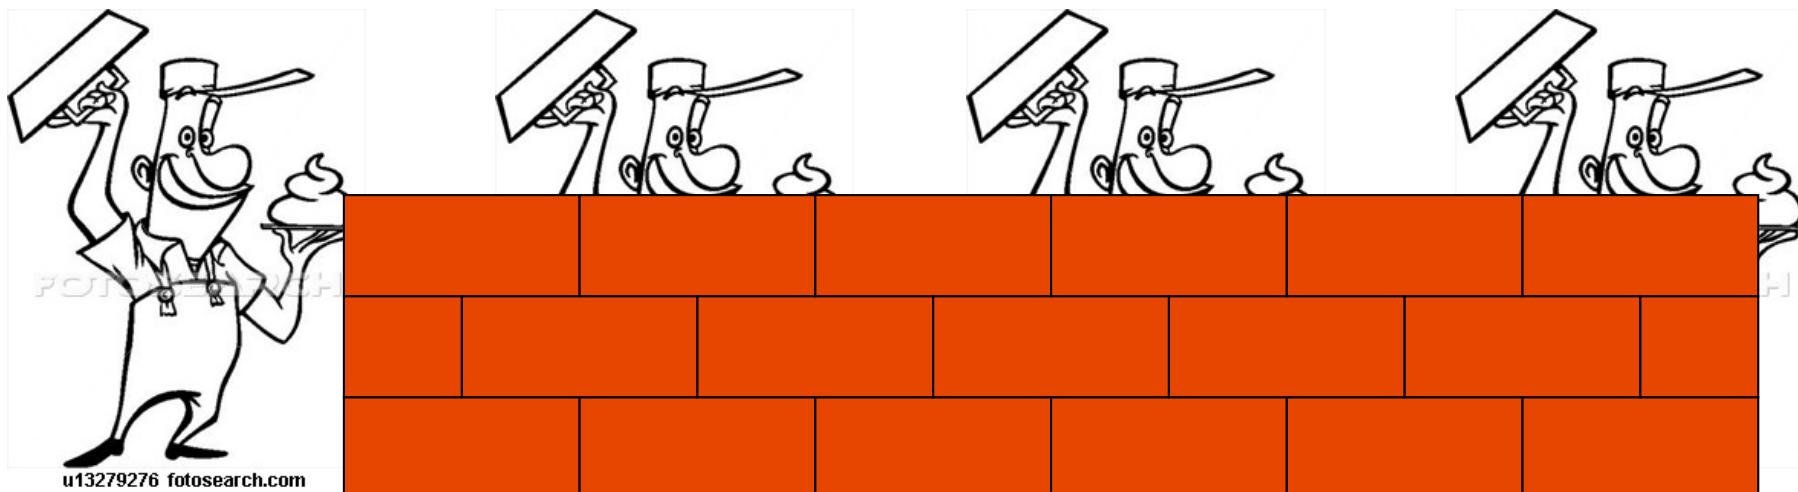
(with apologies to Ian Watson, Univ. Manchester)

PARAPHRASE

How to build a wall faster



How NOT to build a wall

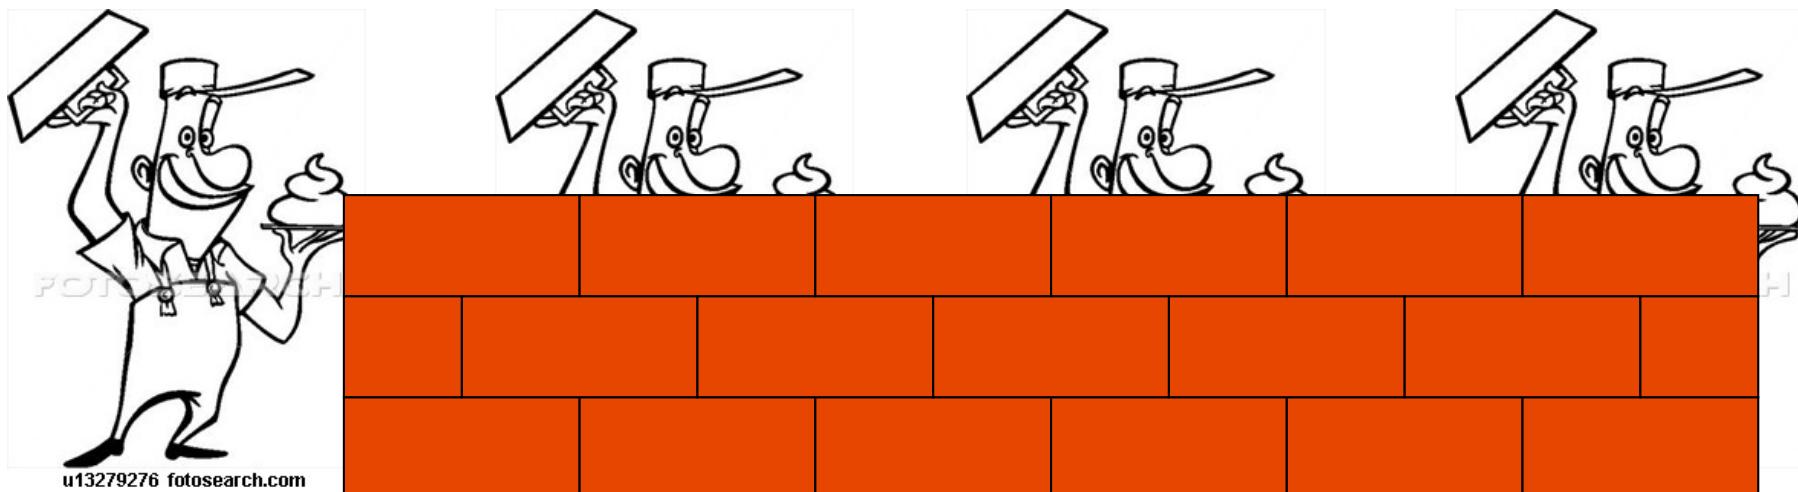


Task identification is not the only problem...

Must also consider Coordination, communication, placement, scheduling, ...

PARAPHRASE

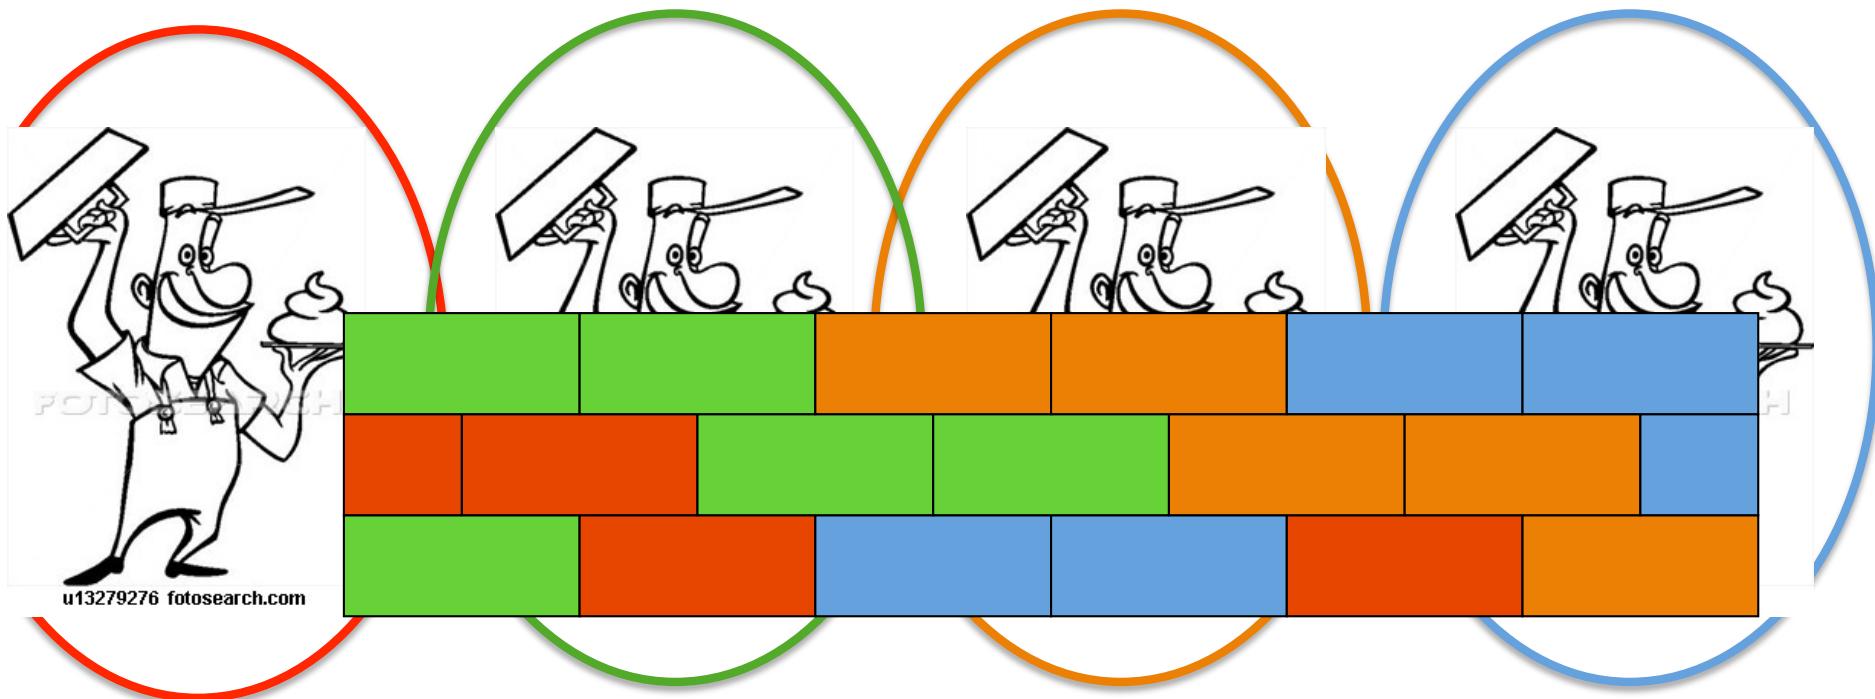
Scheduling



**First row of bricks has to be built before the second
Second row of bricks has to be built before the third**

PARAPHRASE

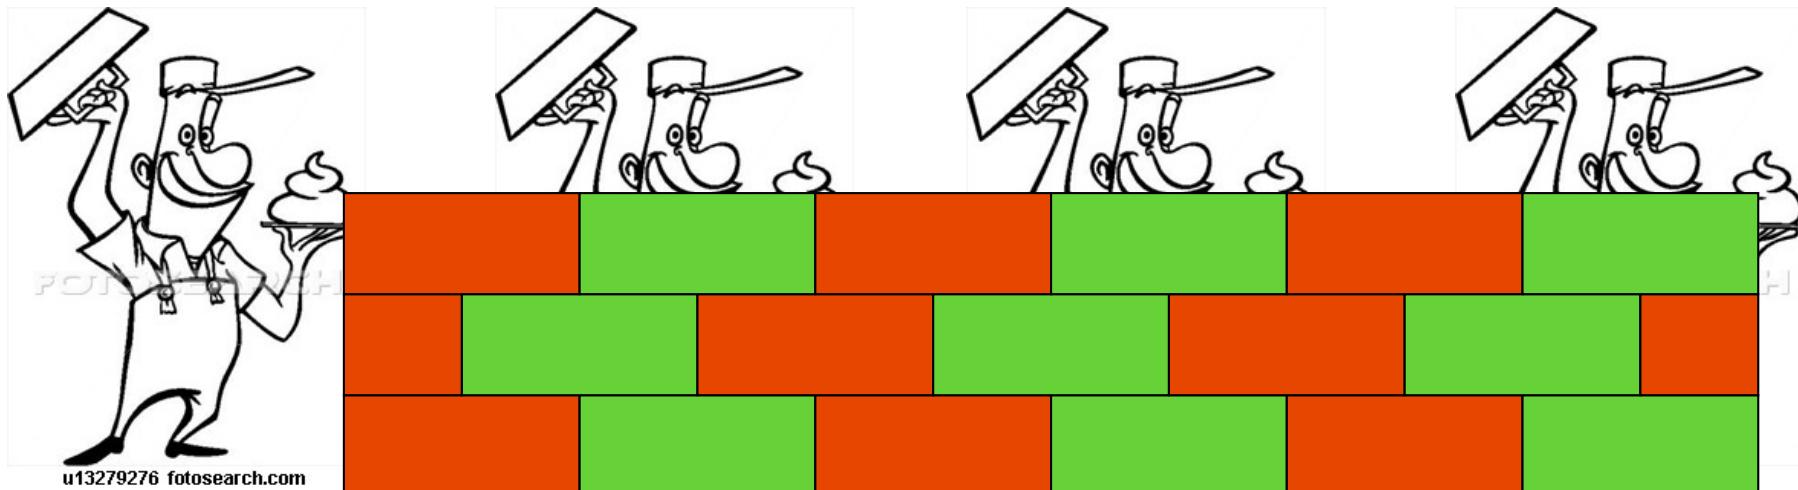
Placement



Which builder places which brick?

PARAPHRASE

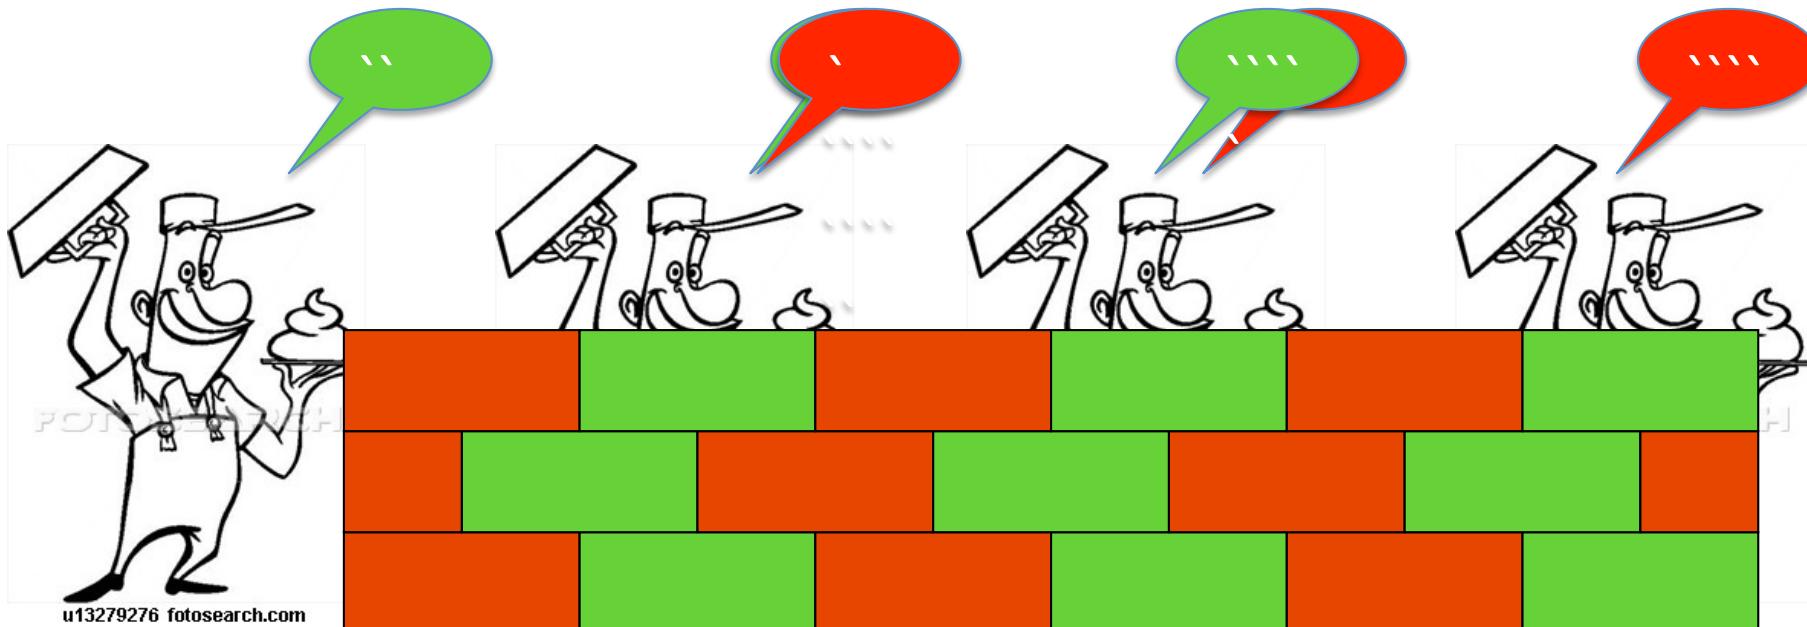
Coordination



Two adjacent bricks cannot be built at the same time

PARAPHRASE

Communication



All builders need to agree when to build a brick...
A-E, F-J, K-N, O-Z

PARAPHRASE

We need structure
We need abstraction

We don't need another brick in the wall

Thinking in Parallel

- Fundamentally, programmers must learn to “think parallel”
 - this requires new *high-level* programming constructs
 - you cannot program effectively while worrying about deadlocks etc
 - they must be eliminated from the design!
 - you cannot program effectively while fiddling with communication etc
 - this needs to be packaged/abstracted!
- Concurrency primitives are NOT the answer!

A Solution?

“The only thing that works for parallelism is functional programming”

Bob Harper, Carnegie Mellon

PARAPHRASE

Parallel Functional Programming

- No explicit ordering of expressions
- Purity means no side-effects
 - Impossible for parallel processes to interfere with each other
 - Can debug sequentially but run in parallel
 - **Enormous** saving in effort
- Programmer concentrate on solving the problem
 - Not porting a sequential algorithm into a (ill-defined) parallel domain
- **No locks, deadlocks or race conditions!!**

PARAPHRASE

Glorious Parallel Haskell (GpH)

GpH aims for an approach where the runtime system manages parallelism which is introduced by the parallel programmer

- Based on the functional language Haskell
- Simple parallel language extension
- Automatic control of parallelism

<http://hackage.haskell.org/platform/>

The *par* Annotation

The basic way to introduce parallelism in GpH. A variant of “lazy thread creation” (Peyton Jones & Salkild, 1989):

a `par` b

create a **spark** for **a** and return the value of **b**

For example

x `par` f x where x = ...

sparks **x**, and returns **f x**

Simple Example: Parallel Fibonacci

```
import Control.Parallel -- for par

pfib :: Int -> Int
pfib n
| n <= 1      = 1
| otherwise = n2 `par` (n1 `par` (n1+n2+1))
  where
    n1 = pfib (n-1)
    n2 = pfib (n-2)
```

Not a real problem, but
easy to understand!

-- seq This includes *all* locking, placement/mapping,
fib ::
fib n
| n -> creation, data/task serialisation, load balancing,
| otherwise -> thread migration, scheduling, parallel garbage
 collection, remote caching, and other memory
 management

Compiling and Running Parallel Haskell

- Use the `-threaded` flag to GHC to enable parallelism

```
$ ghc -threaded -rtsopts --make MyProg.hs
```

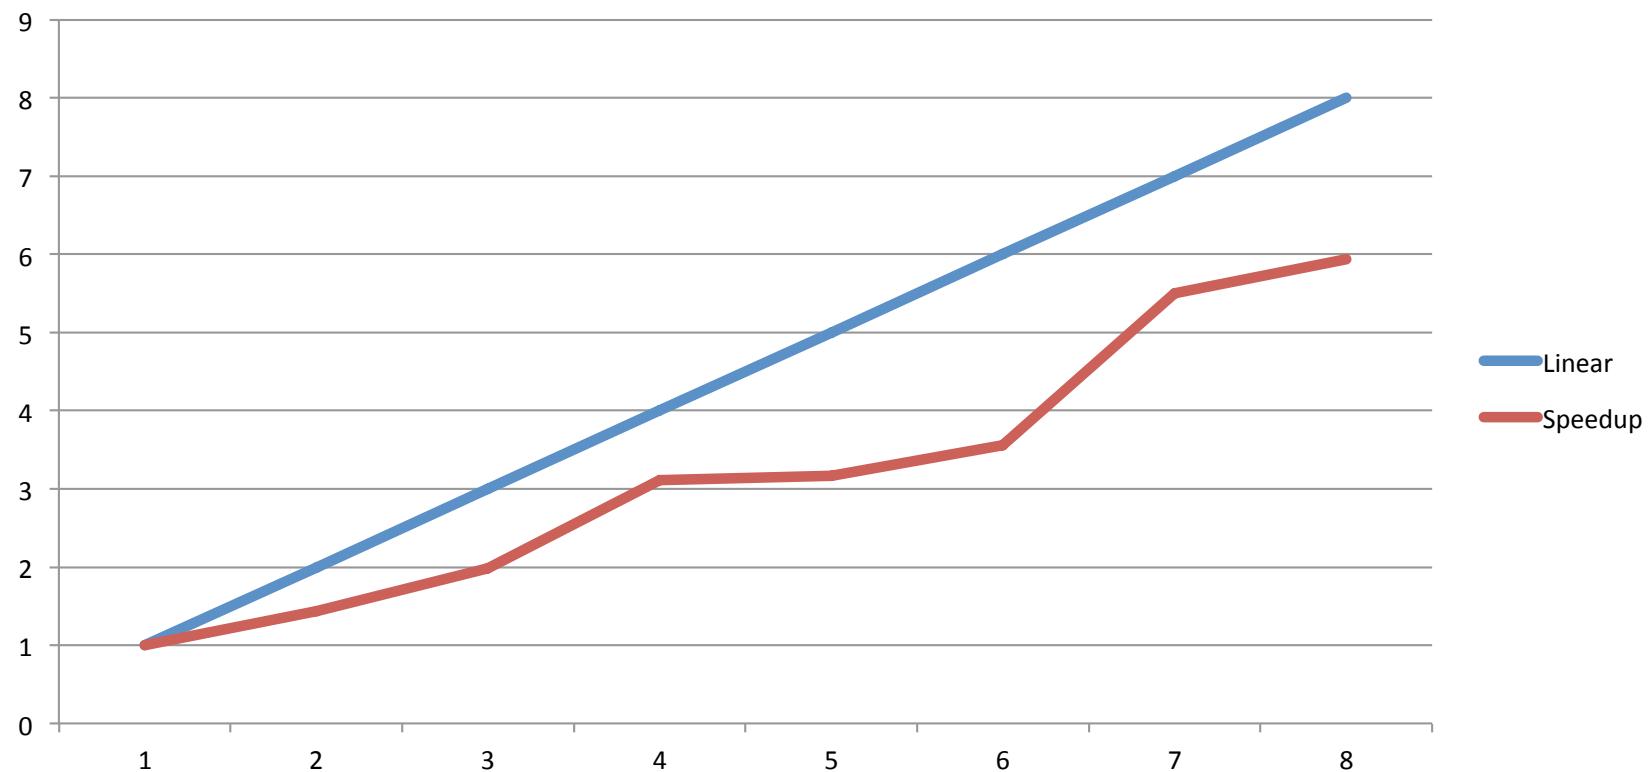
- Run the program using the `-N` option to specify the number of cores

```
$ MyProg +RTS -N16
```

1973

```
$
```

Speedups for Fibonacci



Another Example

- Consider a definition of quicksort in Haskell

```
quicksort [] = []
```

```
quicksort [x] = [x]
```

```
quicksort (x:xs) = losort ++ (x:hisort)
```

where losort = quicksort [y | y <- xs, y < x]

hisort = quicksort [y | y <- xs, y >= x]

- This can be parallelised by spawning off each sub-sort

```
quicksort [] = []
```

```
quicksort [x] = [x]
```

```
quicksort (x:xs) = losort `par` hisort `par` (losort ++ (x:hisort))
```

where losort = quicksort [y | y <- xs, y < x]

hisort = quicksort [y | y <- xs, y >= x]

Sorting Example (2)

- There are two trivial (base) cases

`qsort [] = []`

`qsort [x] = [x]`

- Otherwise, the two sub-parts are *lосort* and *hісort*

`lосort = qsort [y | y <- xs, y < x]`

`hісort = qsort [y | y <- xs, y >= x]`

- These are sparked in parallel

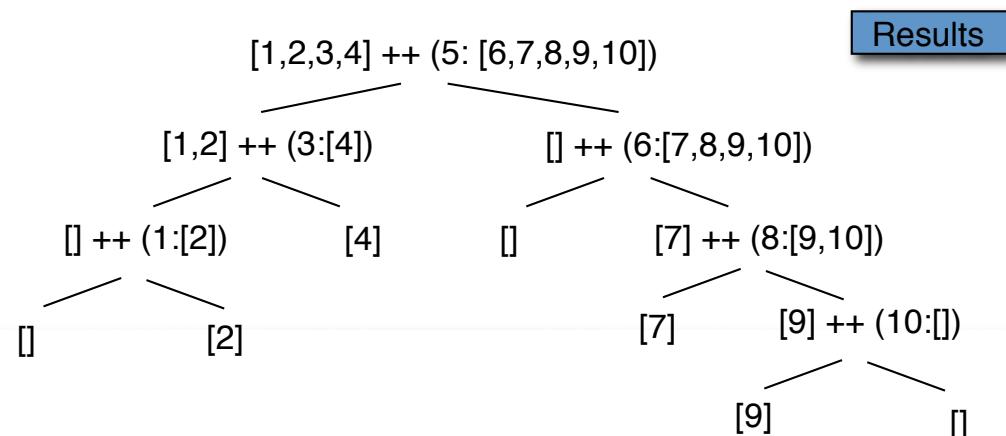
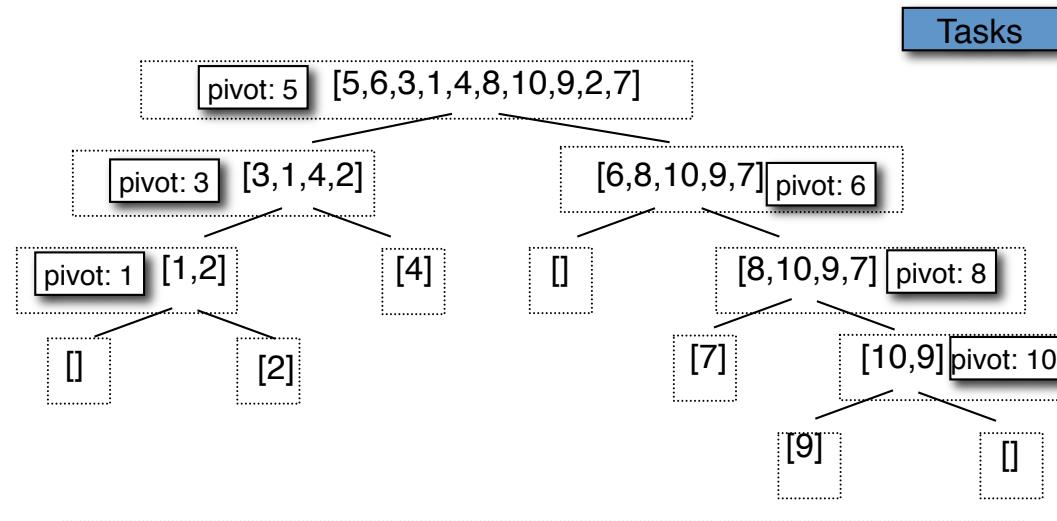
`lосort `par` hісort `par` ...`

- Results from each sub-part are combined using `++` and :

`(lосort ++ (x:hісort))`

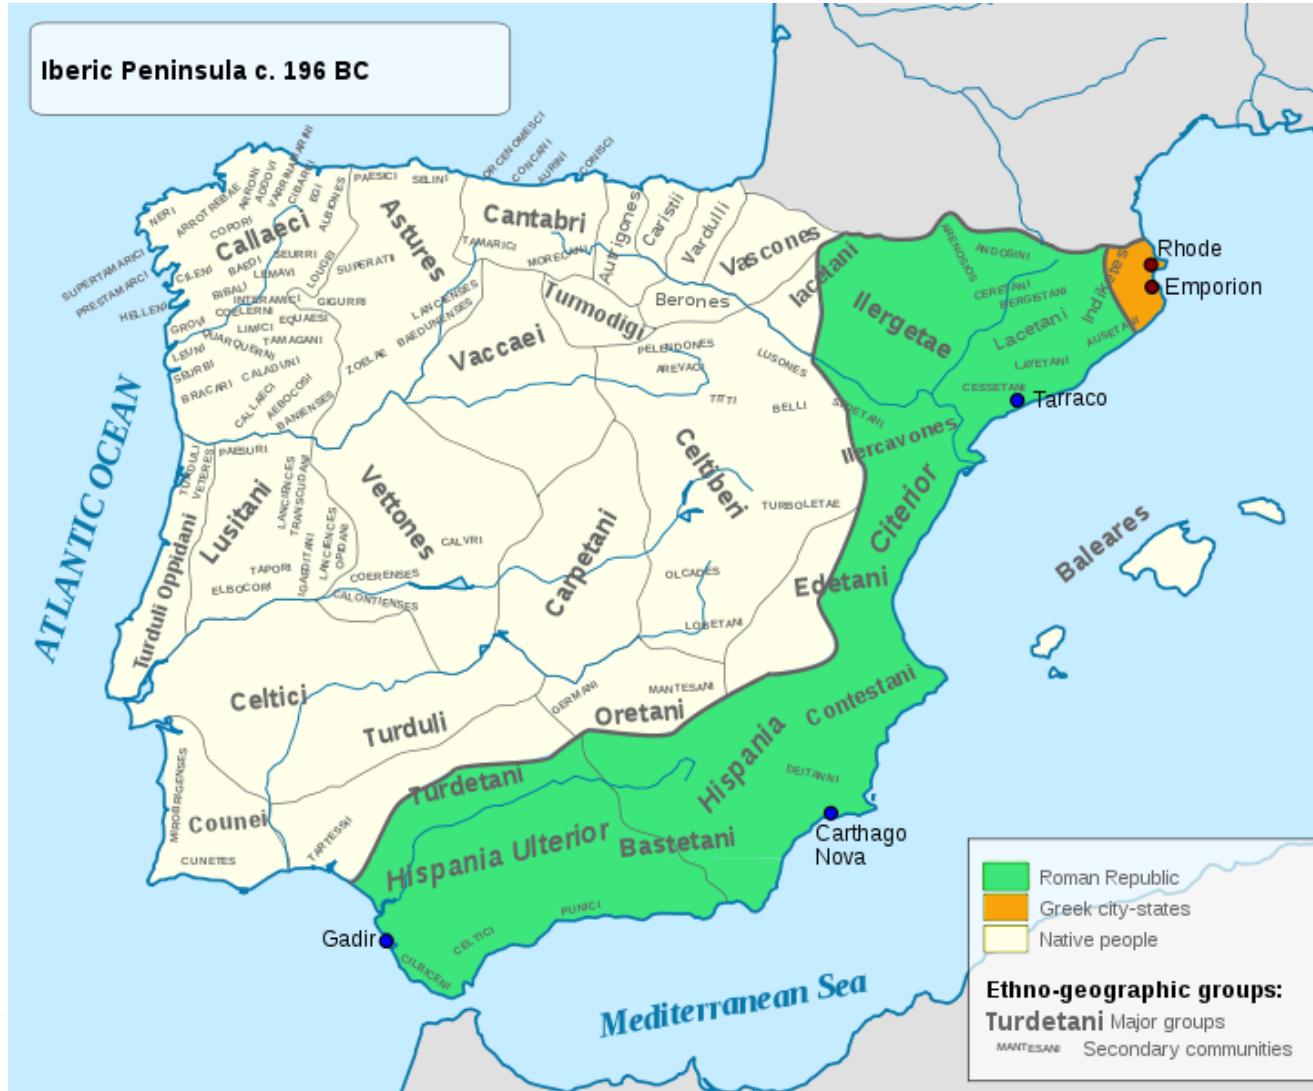


Quicksort [5,6,3,1,8,10,9,2,7]



There's a well-known pattern here

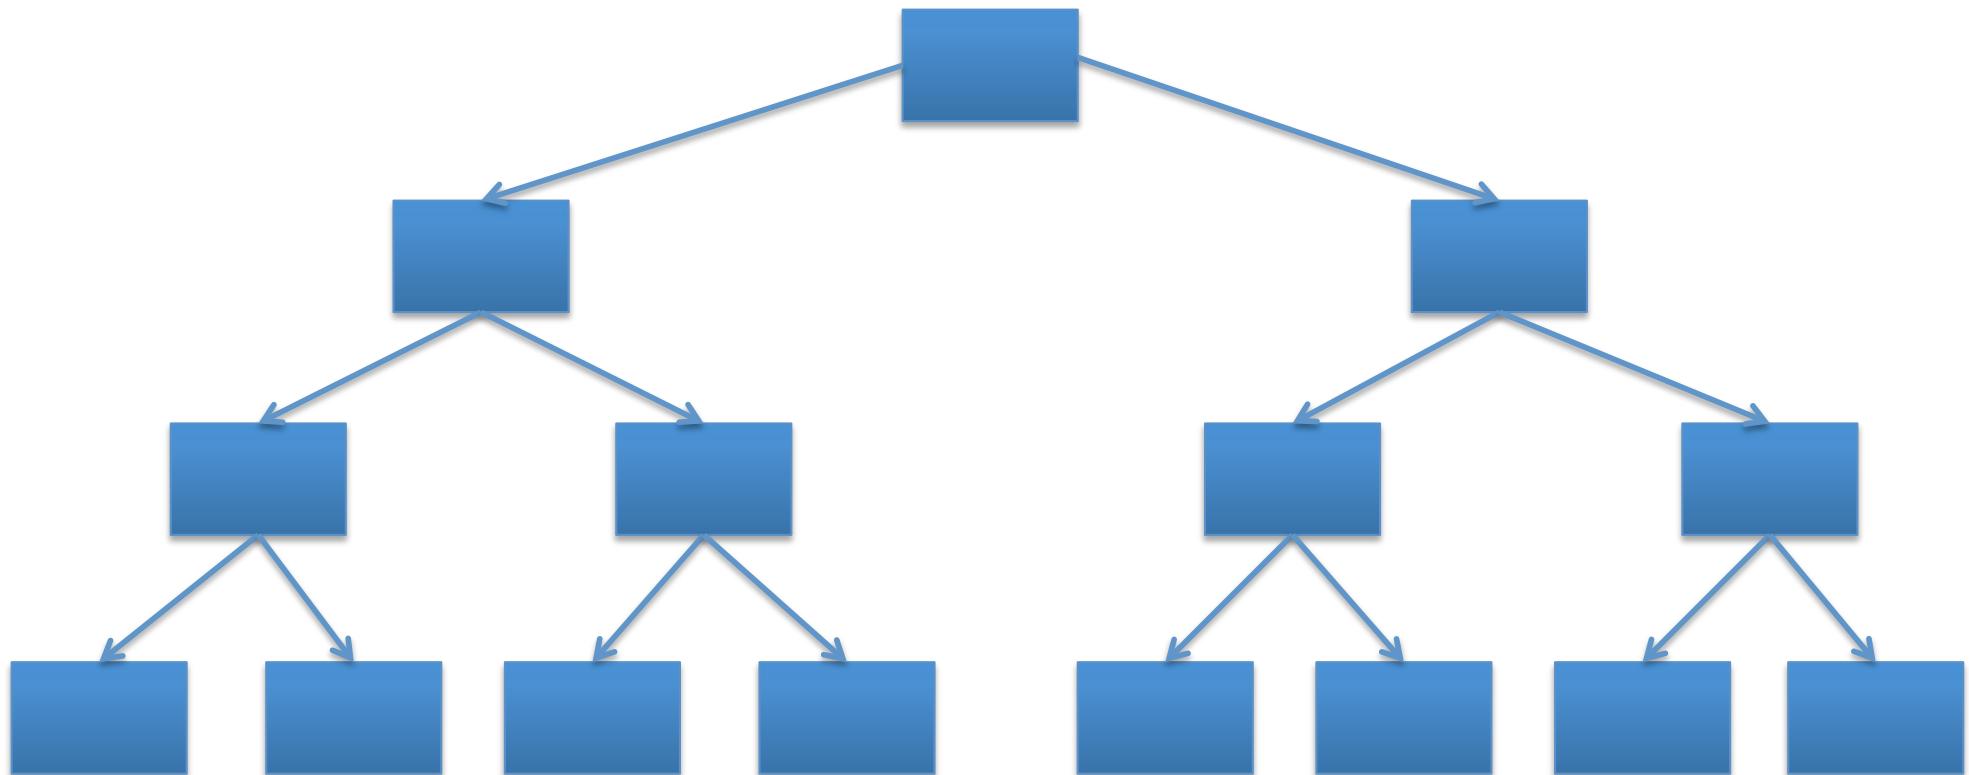
- Divide-and-Conquer
(source: wikipedia)



Divide-and-Conquer

- *Divide* a problem into subparts
- *Solve* each of those parts independently
 - using a divide-and-conquer approach
- When the problem is small enough to be trivially solvable
 - just solve it!!
- *Combine* the results for each solved part

Binary Divide and Conquer



PARAPHRASE

Advantages of DC Parallelism

1. a large number of tasks can be generated very rapidly
2. can be quickly disseminated across a parallel system
 - this is especially important for a distributed-memory system
3. the communication paths are well-defined and hierarchical
 - again, this is important in a distributed-memory system
4. no one processor becomes overloaded with the task of distributing work and marshalling results
 - the work is distributed among the available processors

Parallel Quicksort Revisited

- We parallelised quicksort by spawning off each sub-sort

`qsort [] = []`

`qsort [x] = [x]`

`qsort (x:xs) = losort `par` hisort `par` (losort ++ (x:hisort))`

where `losort = qsort [y | y <- xs, y < x]`

`hisort = qsort [y | y <- xs, y >= x]`

- This is an example of divide-and-conquer parallelism



A general divide-and-conquer Function

- We can write a *general* function to capture divide-and-conquer

`dc split threshold combine worker input = combine results`

where `results =`

```
if threshold input then
    worker input
else parmap
    (dc split threshold combine worker)
    (split input)
```

```
dc :: (a -> [a]) ->          -- split
      (a -> Bool) ->         -- threshold
      ([b] -> b) ->          -- combine
      (a -> [b]) ->          -- worker
      a -> b
```

Skeletons

- *dc* is an example of an *algorithmic skeleton*
 - an *implementation* of a parallel pattern
- Skeletons are templates
 - pluggable *higher-order* functions
 - can be instantiated with concrete worker functions

Murray Cole,

"Algorithmic Skeletons: structured management of parallel computation" MIT Press, 1989

Horacio González-Vélez and Mario Leyton:

"A survey of algorithmic skeleton frameworks: high-level structured parallel programming enablers."

Softw., Pract. Exper. 40(12): 1135-1160 (2010)

Instantiating the DC Skeleton

- dc requires four arguments
 - one *split* function – how to decompose the input
 - one *threshold* function – when is the input simple enough
 - one *combine* function – how to combine the results
 - one *worker* function – what are we actually trying to do

```
dc split threshold combine worker
```

- all of these are function arguments to *dc*

A General Parallel Quicksort

- We can use this to define quicksort

```

qsortDC [] = []
qsortDC list = dc split threshold combine worker list
    where   worker = qsort

        combine = concat

threshold xs = length xs < 1 
split (x:xs) = [lows, [x], highs]
    where   lows = [ y | y <- list, y < x]
            highs = [ y | y <- list, y >= x ]

```

A General Parallel Fibonacci

- We can also use *dc* to define Fibonacci

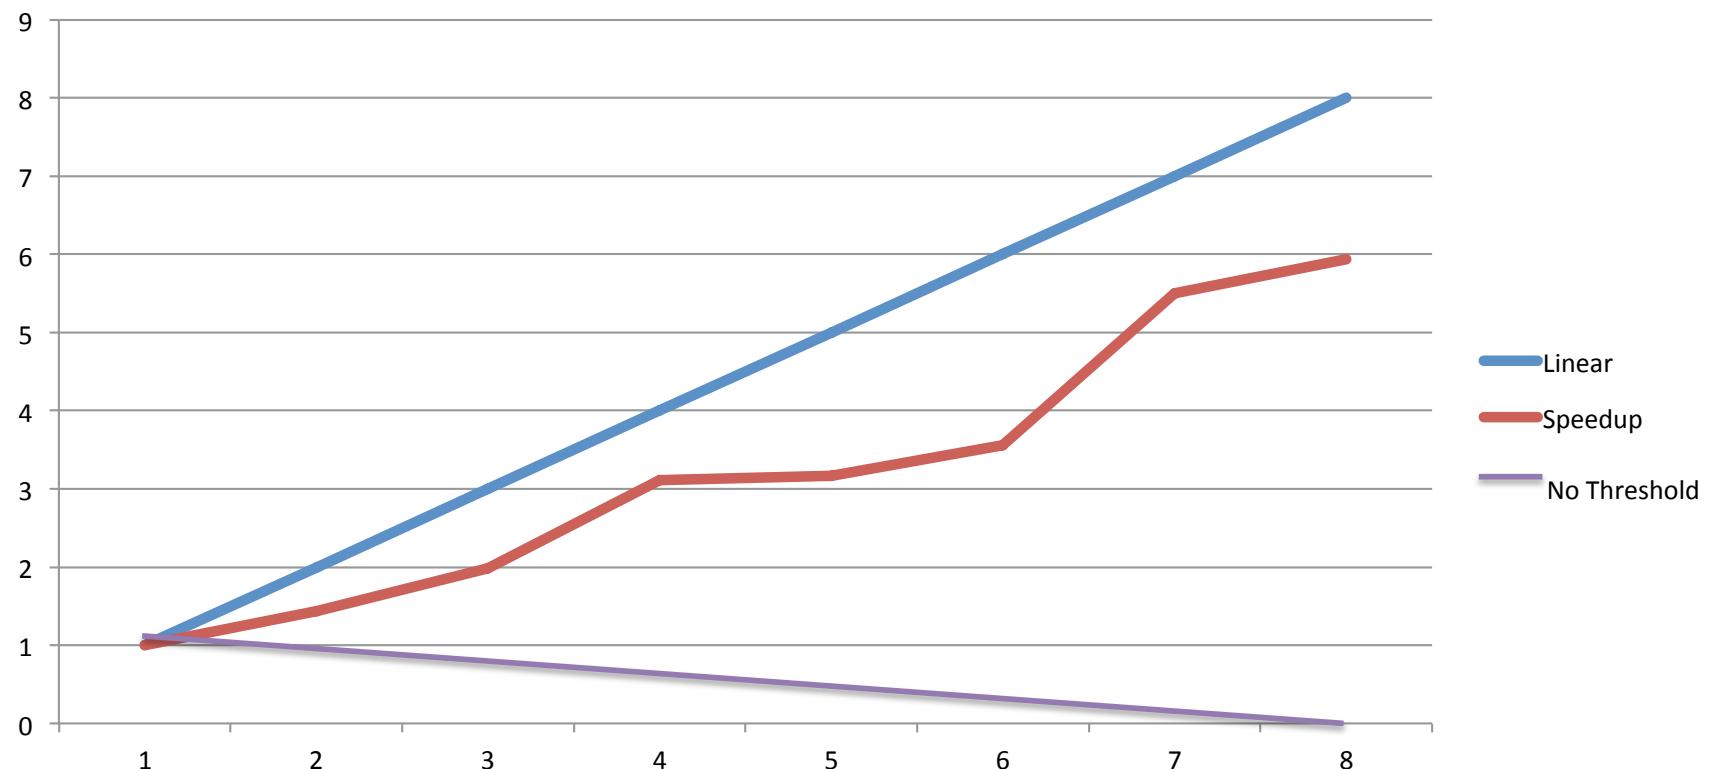
```
fibDC n = dc split threshold combine worker list
          where    worker = fib

          combine = foldr (+) 1

          threshold n = n <= 20

          split n = [n-1,n-2]
```

Speedups for Fibonacci (again)



Parallelism is not Concurrency

- Concurrency is a programming abstraction
 - The *illusion* of independent threads of execution
- Parallelism is a hardware artefact
 - The *reality* of threads executing at the same time
- Concurrency is about breaking a program down into independent units of computation
- Parallelism is about making things happen at the same time
- Typically a parallel program has thousands or millions of threads
- A concurrent program typically has a few threads

Parallelism is not Concurrency (2)

- A concurrent thread may be broken down into many parallel threads
 - or none at all
- Parallelism can sometimes be modelled by concurrency
 - but implicit parallelism cannot!
- Concurrency is about maintaining dependencies
 - Parallelism is about breaking dependencies
- If we try to deal with parallelism using concurrency models/primitives, we are using the wrong abstractions
 - Too low-level, Too coarse-grained, Not scalable

Conclusions

- Functional programming makes it easy to introduce parallelism
 - No side effects means any computation could be parallel
 - Matches pattern-based parallelism
 - Much detail can be abstracted
- Lots of problems can be avoided
 - e.g. Freedom from Deadlock
 - Parallel programs give the same results as sequential ones!
- *But not completely trivial!!*
 - Need to choose granularity carefully!
 - e.g. thresholding
 - May need to understand the execution model
 - e.g. pseq

Isn't this all just wishful thinking?



PARAPHRASE

NO!

- C++11 has lambda
- Java 8 will have lambda (closures)
- Apple uses closures in Grand Central Dispatch

Some Research Challenges

- How do we model parallelism formally
 - when is one program “better” than other
 - Can we prove this???
- How do we deal with the “megacore challenge”
 - scaling, heterogeneity, multiple levels
- What are the best abstractions for parallelism
 - skeletons (what skeletons?), strategies, ...
 - How do we help the programmer “think in parallel”
 - What do we do if a pattern doesn’t quite fit the problem
- How do we understand performance
 - visualisation, abstraction, formal reasoning, ...
- How can we analyse resource usage in parallel systems
 - Time, energy, ...
- What about tool support (e.g. refactoring)
- Can we do it all automatically??

Further Reading

Trinder, Hammond, Loidl and Peyton Jones
“Algorithm + Strategy = Parallelism”,
Journal of Functional Programming, 8(1):23–60, 1998

Brown, Loidl and Hammond
“ParaForming Forming Parallel Haskell Programs using Novel Refactoring Techniques”
To Appear in Proc. 2011 Trends in Functional Programming (TFP), Madrid, Spain, 2012

Ferreiro, Castro, Janjic and Hammond
“Repeating History: Execution Replay for Parallel Haskell Programs”
Submitted to 2012 Trends in Functional Programming (TFP), St Andrews, June 2012



University
of
St Andrews

Research Directions in Parallel Functional Programming eBook: Kevin Hammond, Greg Michaelson: Amazon.co.uk: Kindle Store

http://www.amazon.co.uk/Research-Directions-Functional-Programming-ebook/dp/B000W67U4G/ref=sr_1_fkmr2_1?ie=UTF8&qid=13391411

Click to **LOOK INSIDE!**

Research Directions in Parallel Functional Programming [Kindle Edition]
Kevin Hammond (Author, Editor), Greg Michaelson (Author, Editor)

Buy now with 1-Click®

LOOK INSIDE! Kindle Book Print Book Zoom - | Zoom + Feedback | Help | Expanded View | Close

Just so you know...
This view is of the print book (Paperback edition). A preview of the Kindle book is currently not available.

Paperback £99.00
Add to Basket
1 used & new from £99.00

Book sections

- Front Cover
- Copyright
- Table of Contents
- First Pages
- Index
- Back Cover
- Surprise Me!

Search Inside This Book

Page Numbers Source ISBN: 1852330929

Research Directions in Parallel Functional Programming

Kevin Hammond and Greg Michaelson (Eds)

The cover features a blue background with two large, stylized white Greek letter 'lambda' symbols. A double-headed arrow points between them, indicating parallelism or functional programming concepts.

Your Browsing History
Page 1 of 1

Research Directions...
by Kevin Hammond, Greg...
£72.43
Look Inside This Book

Edit your book history

Customers Also Bought
Books customers also bought could not be retrieved

Hide These Books

PARAPHRASE

Parallel Haskell: Lightweight Parallelism for Heavyweight Functional Programs

(DRAFT – please do not redistribute without permission.)

Kevin Hammond, Chris Brown and Phil Trinder

In Preparation

Pattern Cookbook

Pattern	Instances	Type	When to use
Divide-and-Conquer	dC dc parThreshold	control-parallel	The operation can be defined recursively, subdividing its input at each recursive call; each application of the operation is independent; the results of recursive calls are combined into a final result; arguments may be created dynamically. Divide-and-conquer works well in a high latency (distributed) setting as well as a low latency (shared memory) setting, since there is no communication hotspot through a central master processor, and tasks can be created in a distributed rather than centralised way.

PARAPHRASE

Funded by

- SCIEnce (EU FP6), Grid/Cloud/Multicore coordination
 - €3.2M, 2005-2012
- Advance (EU FP7), Multicore streaming
 - €2.7M, 2010-2013
- HPC-GAP (EPSRC), Legacy system on thousands of cores
 - £1.6M, 2010-2014
- Islay (EPSRC), Real-time FPGA streaming implementation
 - £1.4M, 2008-2011
- ParaPhrase (EU FP7), Patterns for heterogeneous multicore
 - €2.6M, 2011-2014



Industrial Connections

SAP GmbH, Karlsruhe

BAe Systems

Selex Galileo

Biold GmbH, Stuttgart

Philips Healthcare

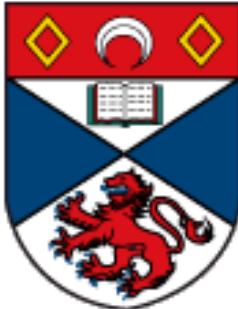
Software Competence Centre, Hagenberg

Mellanox Inc.

Erlang Solutions Ltd

Microsoft Research

Well-Typed



Funded PhD Research Studentship

University of St Andrews - School of Computer Science

The School of Computer Science at the University of St Andrews has funding for students to undertake PhD research in any of the general research areas in the school:

<http://www.cs.st-andrews.ac.uk/research>

We are looking for highly motivated research students with an interest in these exciting research areas. Our only requirements are that the proposed research would be good, we have staff to supervise it, and that you would be good at doing it. We have up to 5 funded studentships available for students interested in working towards a PhD. The studentships offer costs of fees and an annual tax-free maintenance stipend of about £13,590 per year for 3.5 years. Exceptionally well qualified and able students may be awarded an enhanced stipend of an additional £2,000 per year. Students should normally have or expect at least an upper-2nd class Honours degree or Masters degree in Computer Science or a related discipline.

For further information on how to apply, see our postgraduate web pages (<http://www.cs.st-andrews.ac.uk/prospective-pg>). The closing date for applications is **July 20th 2012** and we will make decisions on studentship allocation by **Sep 1st 2012**. (Applications after **July 20th** may be considered, at our discretion.) Informal enquiries can be directed to pg-admin-cs@st-andrews.ac.uk or to potential supervisors.

1	1	11	18	30	1E
2	2	12	19	31	1E
3	3	13	20	32	1E

1	1	11	18	30	1E
2	2	12	19	31	1E
3	3	13	20	32	1E

PARAPHRASE

1	1	11	18	30	1E
2	2	12	19	31	1E
3	3	13	20	32	1E

1	1	11	18	30	1E
2	2	12	19	31	1E
3	3	13	20	32	1E

THANK YOU!

<http://www.paraphrase-ict.eu>

<http://www.project-advance.eu>

@paraphrase_fp7